



exact_buf_write/256 KiB



Additional Statistics:

| | Lower bound | Estimate | Upper bound |
|----------------|--------------|--------------|--------------|
| Slope | 71.903 ns | 72.184 ns | 72.523 ns |
| Throughput | 3366.4 GiB/s | 3382.2 GiB/s | 3395.4 GiB/s |
| R ² | 0.9411677 | 0.9431364 | 0.9402810 |
| Mean | 72.319 ns | 72.678 ns | 73.045 ns |
| Std. Dev. | 1.6104 ns | 1.8683 ns | 2.1019 ns |
| Median | 71.748 ns | 71.839 ns | 72.736 ns |
| MAD | 895.93 ps | 1.4998 ns | 2.1414 ns |

Additional Plots:

- Typical
- Mean
- Std. Dev.
- Median
- MAD
- Slope



Understanding this report:

The plot on the left displays the average time per iteration for this benchmark. The shaded region shows the estimated probability of an iteration taking a certain amount of time, while the line shows the mean. Click on the plot for a larger view showing the outliers.

The plot on the right shows the linear regression calculated from the measurements. Each point represents a sample, though here it shows the total time for the sample rather than time per iteration. The line is the line of best fit for these measurements.

See [the documentation](#) for more details on the additional statistics.

Change Since Previous Benchmark



Additional Statistics:

| | Lower bound | Estimate | Upper bound |
|----------------------|-------------|----------|-------------|
| Change in time | -37.945% | -37.559% | -37.162% |
| Change in throughput | +61.147% | +60.151% | +59.140% |

Performance has improved. (p = 0.00 < 0.05)

Additional Plots:

- Change in mean
- Change in median
- T-Test

Understanding this report:

The plot on the left shows the probability of the function taking a certain amount of time. The red curve represents the saved measurements from the last time this benchmark was run, while the blue curve shows the measurements from this run. The lines represent the mean time per iteration. Click on the plot for a larger view.

The plot on the right shows the two regressions. Again, the red line represents the previous measurement while the blue line shows the current measurement.

See [the documentation](#) for more details on the additional statistics.

This report was generated by Criterion.rs, a statistics-driven benchmarking library in Rust.